Neural Networks

0. Logistics

Spring 2019

Neural Networks are taking over!

- Neural networks have become one of the major thrust areas recently in various pattern recognition, prediction, and analysis problems

- In many problems they have established the state of the art
  - Often exceeding previous benchmarks by large margins

Breakthroughs with neural networks

TECHNEWSWORLD EMERGING TECH SEARCH

Computing Internet IT Mobile Tech Reviews Security Technology Tech Blog Reader Services

Microsoft AI Beats Humans at Speech Recognition

By Richard Adhikari
Oct 20, 2016 11:40 AM PT

Print Email

G+1 5 Tweet 25 Share 45 Share 11 Share 0 share 104

Image: Adobe Stock

Microsoft's Artificial Intelligence and Research Unit earlier this week reported that its speech recognition technology had surpassed the performance of human transcriptionists.

Most Popular Newsletters News Alerts

How do you feel about Black Friday and Cyber Monday?

- They're great -- I get a lot of bargains!
- The deals are too spread out -- I'd prefer just one day.
- They're a fun way to kick off the holiday season.
- I don't like the commercialization of Thanksgiving Day.
- They're crucial for the retail industry and the economy.
- The deals typically aren't that good.

Vote to See Results

E-Commerce Times

- Black Friday Shoppers Hungry for New Experiences, New Tech
- Pay TV's Newest Innovation: Giving Users Control
- Apple Celebrates Itself in $300 Coffee Table Tome
- AWS Enjoys Top Perch in IaaS, PaaS Markets
- US Comptroller Gears Up for Blockchain and

Breakthroughs with neural networks

Image segmentation & recognition

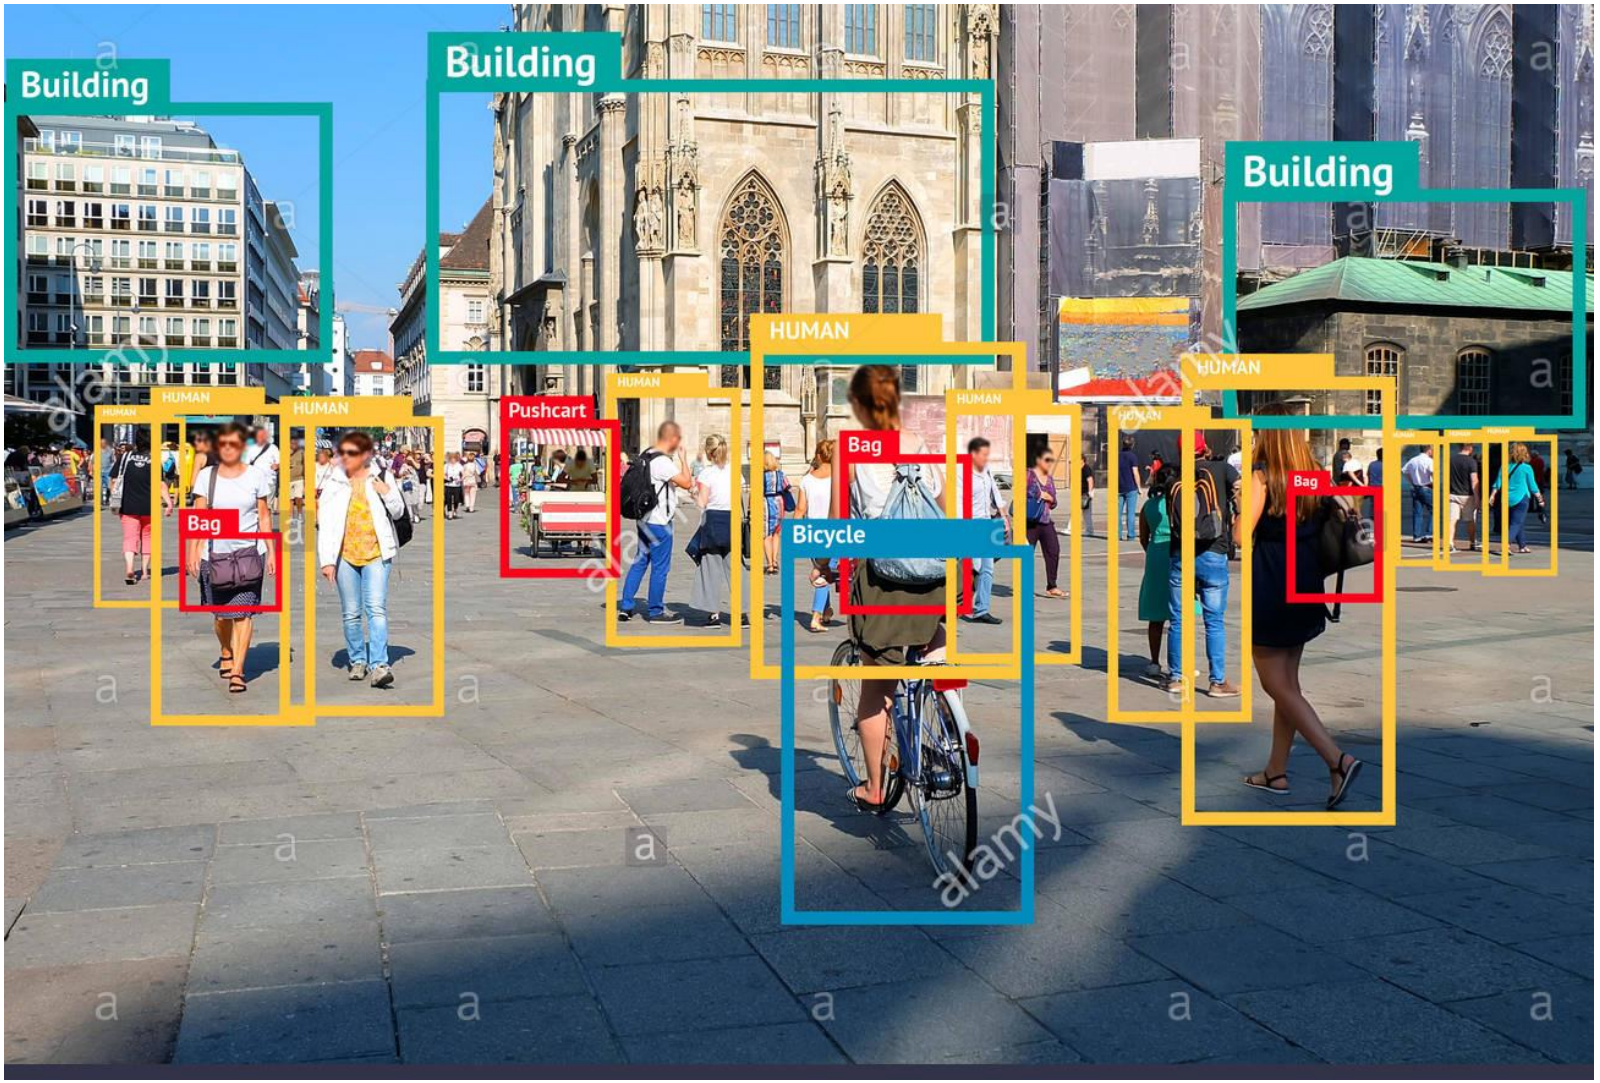

Image recognition

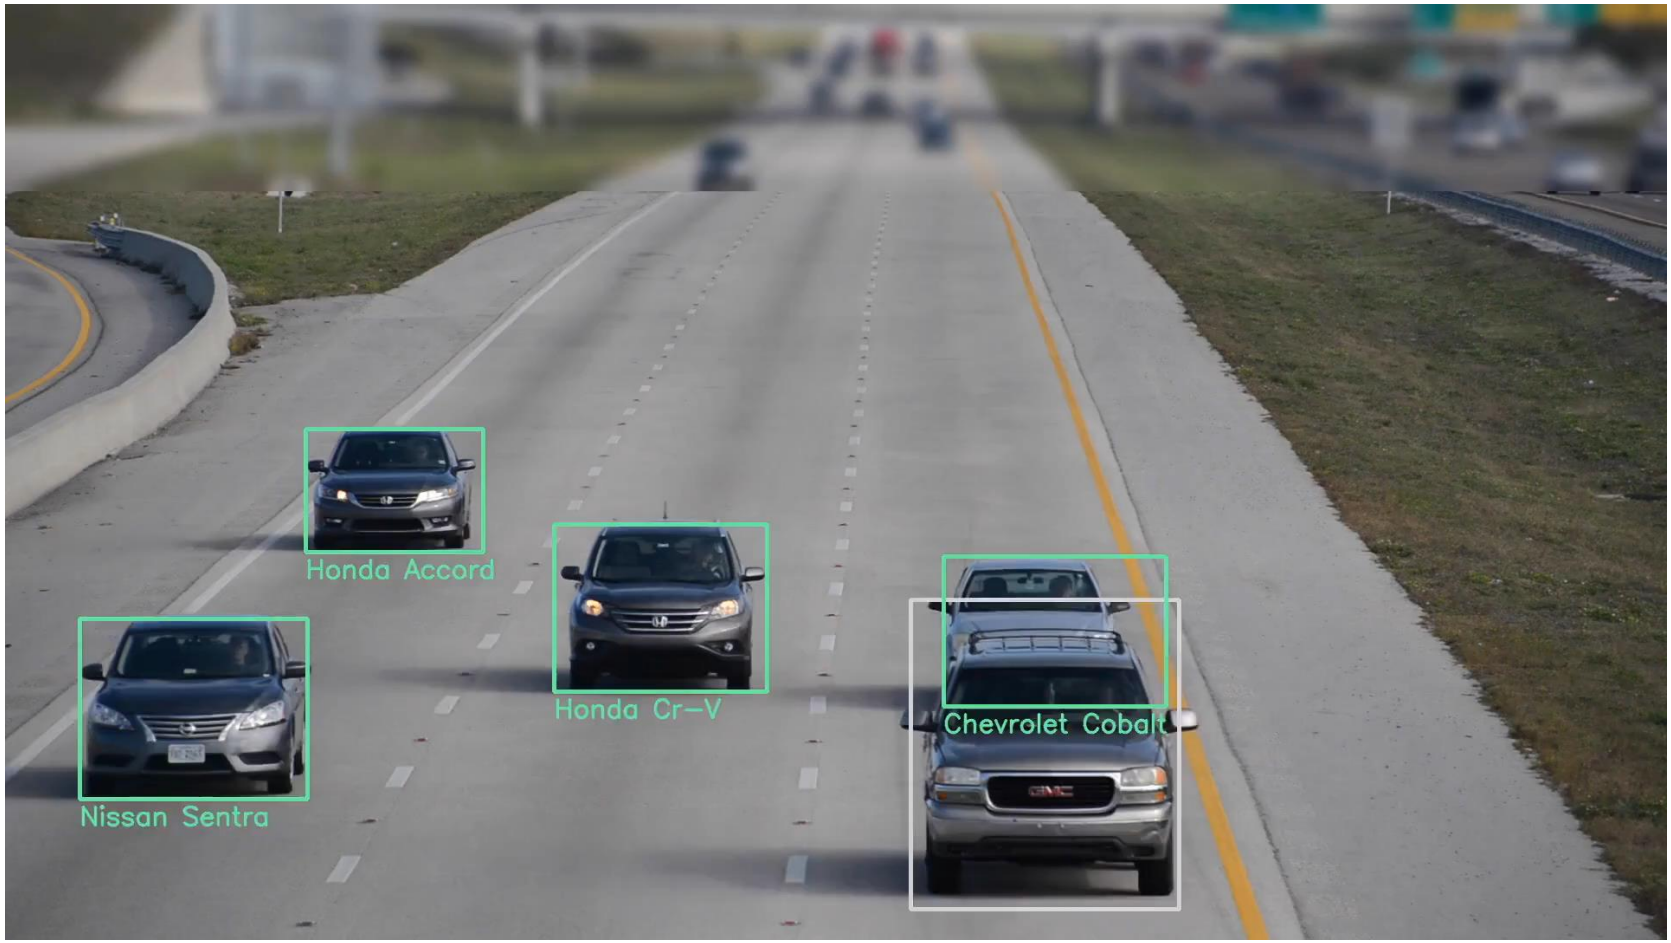

https://www.sighthound.com/technology/

Breakthroughs with neural networks

AlphaGo

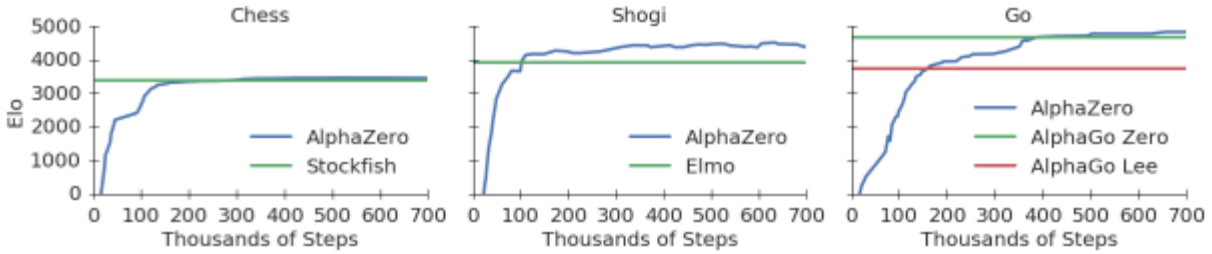

Figure 1: Training *AlphaZero* for 700,000 steps. Elo ratings were computed from evaluation games between different players when given one second per move. **a** Performance of *AlphaZero* in chess, compared to 2016 TCEC world-champion program *Stockfish*. **b** Performance of *AlphaZero* in shogi, compared to 2017 CSA world-champion program *Elmo*. **c** Performance of *AlphaZero* in Go, compared to *AlphaGo Lee* and *AlphaGo Zero* (20 block / 3 day) (*29*).

Breakthroughs with neural networks

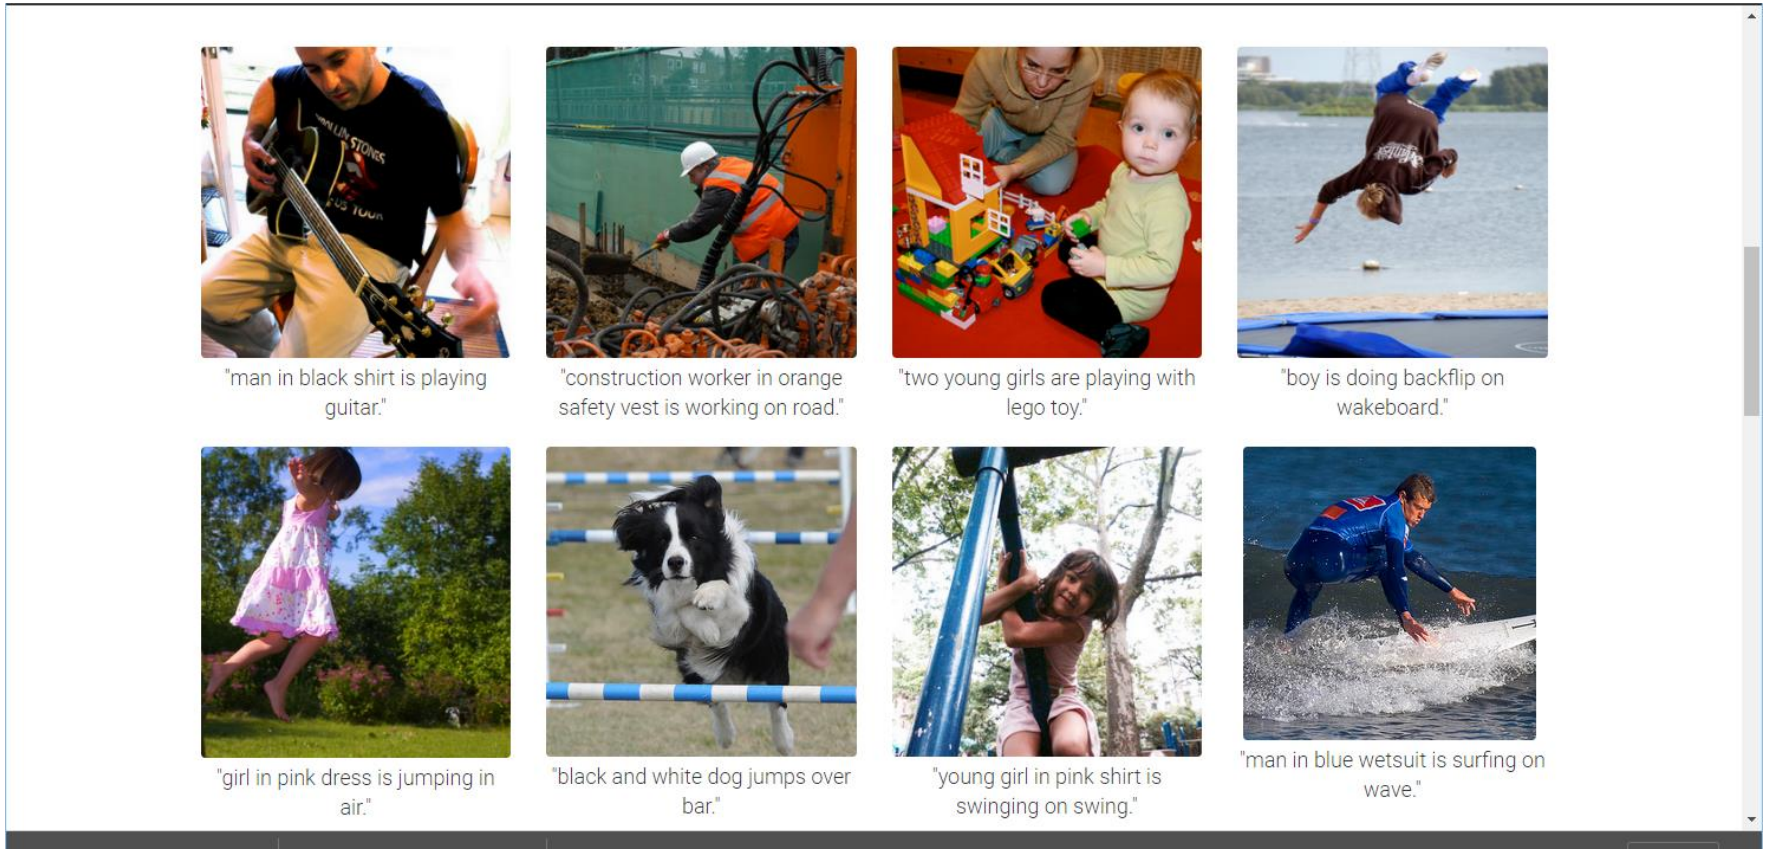

- Captions generated entirely by a neural network

Successes with neural networks

- And a variety of other problems:
  - From art to astronomy to healthcare...
  - and even predicting stock markets!

Neural Networks and the Job Market

This guy didn't know about neural networks (a.k.a deep learning)

This guy learned about neural networks (a.k.a deep learning)

Course Objectives

- Understanding neural networks
- Comprehending the models that do the previously mentioned tasks
  - And maybe build them
- Familiarity with some of the terminology
  - What are these:
    - http://www.datasciencecentral.com/profiles/blogs/concise-visual-summary-of-deep-learning-architectures
- Fearlessly design, build and train networks for various tasks

- *You will not become an expert in one course*

Course objectives: Broad level

- Concepts
  - Some historical perspective
  - Types of neural networks and underlying ideas
  - Learning in neural networks
    - Training, concepts, practical issues
  - Architectures and applications
  - Will try to maintain balance between squiggles and concepts (concept >> squiggle)
- Practical
  - Familiarity with training
  - Implement various neural network architectures
  - Implement state-of-art solutions for some problems
- Overall: Set you up for further research/work in your research area

Course learning objectives: Topics

- Basic network formalisms:
  - MLPs
  - Convolutional networks
  - Recurrent networks
  - Boltzmann machines
- Some advanced formalisms
  - Generative models: VAEs
  - Adversarial models: GANs
- Topics we will touch upon:
  - Computer vision: recognizing images
  - Text processing: modelling and generating language
  - Machine translation: Sequence to sequence modelling
  - Modelling distributions and generating data
  - Reinforcement learning and games
  - Speech recognition

Reading

- List of books on course webpage
- Additional reading material also on course pages

Instructors and TAs

- Instructor: Bhiksha Raj
  - bhiksha@cs.cmu.edu
  - x8-9826
- TAs:
  - List of TAs, with email ids on course page
  - We have TAs for the
    - Pitt Campus
    - Kigali,
    - SV campus,
    - Doha campus
  - Please approach your local TA first
- Office hours: On webpage
- http://deeplearning.cs.cmu.edu/

Logistics: Lectures..

- Have in-class and online sections
  - Including online sections in Kigali, SV and Doha
- Lectures are streamed
- Recordings will be posted
- Important that you view the lectures
  - Even if you think you know the topic
  - Your marks depend on viewing lectures

Lecture Schedule

- On website
  - The schedule for the latter half of the semester may vary a bit
    - Guest lecturer schedules are fuzzy..
- Guest lectures:
  - TBD
    - Scott Fahlman, Graham Neubig, etc.

Recitations

- We will have 13 recitations
- Will cover implementation details and basic exercises
  - Very important if you wish to get the maximum out of the course
- Topic list on the course schedule

- *Strongly recommend attending all recitations*
  - *Even if you think you know everything*

Recitations Schedule

- 16 Jan 2019 AWS
- 25 Jan 2019 Your first Deep Learning Code
- 1 Feb 2019 Efficient Deep Learning/Optimization Methods
- 8 Feb 2019 Debugging and Visualization
- 15 Feb 2019 Convolutional Neural Networks
- 22 Feb 2019 CNNs: HW2
- 1 Mar 2019 Recurrent Neural Networks
- 8 Mar 2019 RNN: CTC
- 22 Mar 2019 Attention
- 29 Mar 2019 Variation Auto Encoders
- 5 Apr 2019 GANs
- 19 Apr 2019 Boltzmann machines
- 26 Apr 2019 Reinforcement Learning

See course page for exact details!

Grading

Weekly Quizzes		**24%**
14 Quizzes, bottom two dropped		24%
Assignments		**51%**
HW0 – Preparatory homework	(AL)	1%
HW1 – Basic MLPs	(AL + Kaggle)	12.5%
HW2 – CNNs	(AL + Kaggle)	12.5%
HW3 – RNNs	(AL + Kaggle)	12.5%
HW4 – Sequence to Sequence Modelling	(Kaggle)	12.5%
Team Project (11-785 only)		**25%**
Proposal		TBD
Mid-term Report		TBD
Project Presentation		TBD
Final report		TBD

Weekly Quizzes

- 10-12 multiple-choice questions
- Related to topics covered that week
  - On both slides and in lecture
- Released Friday, closed Saturday night
  - This may occasionally shift, don't panic!
- There will be 14 total quizzes
  - We will consider the best 12
  - This is expected to account for any circumstance-based inability to work on quizzes
    - You could skip up to 2

Lectures and Quizzes

- Slides often contain a lot more information than is presented in class

- *Quizzes will contain questions from topics that are on the slides, but not presented in class*

- *Will also include topics covered in class, but not on online slides!*

Homeworks

- Homeworks come in two flavors
  - Autograded homeworks with deterministic solutions
    - You must upload them to autolab
  - Kaggle problems
    - You compete with your classmates on a leaderboard
    - **We post performance cutoffs for A, B and C**
      - **If you achieved the posted performance for, say "B", you will at least get a B**
      - **A+ == 105 points (bonus)**
      - **A = 100**
      - **B = 80**
      - **C = 60**
      - **D = 40**
      - **No submission: 0**
    - **Actual scores are linearly interpolated between grade cutoffs**
      - Interpolation curves will depend on distribution of scores

Homework Deadlines

- Multiple deadlines
- Separate deadline for Autograded deterministic component
- Kaggle component has multiple deadlines
  - *Initial submission* deadline : If you don't make this, all subsequent scores are multiplied by 0.9
  - *Full submission deadline:* Your final submission must occur before this deadline to be eligible for full marks
  - *Drop-dead deadline:* Must submit by here to be eligible for *any* marks
    - Day on which solution is released

- Homeworks: Late policy
  - Everyone gets up to 7 total slack days (does not apply to initial submission)
  - You can distribute them as you want across your HWs
    - You become ineligible for "A+" bonus if you're using your grace days for Kaggle
  - Once you use up your slack days, all subsequent late submissions will accrue a 10% penalty (on top of any other penalties)
  - There will be no more submissions after the drop-dead deadline
  - Kaggle: Kaggle leaderboards stop showing updates on full-submission deadline
    - But will continue to privately accept submissions until drop-dead deadline

- ***Please see course webpage for complete set of policies***

Preparation for the course

- Course is implementation heavy
  - A lot of coding and experimenting
  - Will work with some large datasets

- Language of choice: Python
- Toolkit of choice: Pytorch
  - You are welcome to use other languages/toolkits, but the TAs will not be able to help with coding/homework
    - Some support for TensorFlow

- We hope you have gone through
  - Recitation zero
  - HW zero
    - Carries marks

Additional Logistics

- Discussions:
  - On Piazza
- Compute infrastructure:
  - Everyone gets Amazon tokens
  - Initially a token for $50
  - Can get additional tokens of $50 up to a total of $150

This course is not easy

- A lot of work!

This course is not easy

- A lot of work!
- *A lot of work!!*

This course is not easy

- A lot of work!
- *A lot of work!!*
- ***A lot of work!!!***

This course is not easy

- A lot of work!
- *A lot of work!!*
- ***A lot of work!!!***
- ***A LOT OF WORK!!!!***

Not for chickens!

This course is not easy

- A lot of work!
- *A lot of work!!*
- ***A lot of work!!!***
- ***A LOT OF WORK!!!!***
- *Mastery-based* evaluation
  - Quizzes to test your understanding of topics covered in the lectures
  - HWs to teach you to implement complex networks
    - And optimize them to high degree
- Target: Anyone who gets an "A" in the course is technically ready for a deep learning job

HW0 / Recitation 0

- Please, please, please, please, please go through the videos for recitation 0, and complete HW0.
  - These are essential for you to gain comfort with the coding require in the following homeworks

- HW1 part 1 also has many components intended to help you *later* in the course
  - So if it seems a bit dense, please bear with it, its worth it
- HW1 is the easiest HW!

Questions?

- Please post on piazza